



Reasoning vs Instruct Models: Performance and Memory Comparison


**Llama-8B Performance
(Green=Better)**



**Qwen-7B Performance
(Green=Better)**



**Llama-8B Memory
(Green=Savings)**



**Qwen-7B Memory
(Green=Savings)**

